

1. What data structures do you use for storing the inverted index? Do you use skip pointers? Explain your choice. In which format do you store the index on disk?

I used a nested dictionary structure (dictionary in a dictionary). A python dictionary is a hashmap, with worst-case  $O(n)$  and average time complexity  $O(1)$ . The dictionary provides an efficient way to read from memory and access positing lists corresponding to specific words in the corpus. The outer dictionary key represents a stemmed word, while the inner dictionary holds the document ID and a list of word positions in the document. The inverted index was stored in a JSON file format, because it provided me with a highly structured data file that I could easily read in my query.py program.

2. How would your data structures for postings change if new documents were added frequently to the collection?

If I were adding new documents often, I would consider adding a linked list structure. Linked-lists are very efficient ( $O(1)$ ) for adding new items to the end of a sequence, since the next pointer is set to this last item, with a reference typically stored to the previous last item. In contrast, adding many new items to an array can be inefficient ( $O(n)$ ), since if the array is already full, it must be copied to add a new element.

3. Examine the lengths of the postings lists for the terms in the full collection and comment on their distribution

I would compute mean and variance of the posting lists corresponding to the terms. I ran out of time to complete this analysis.

4. If you were not provided with the stop words list, how would you have created one?

If not provided with a stop words list, I would consider using the term frequency-inverse document frequency (tf-idf) statistic to help me generate a stop-word list for a given set of documents. I would want to compute the amount of times a word appears in a given document. I would search for words with a low tf-idf score, because these are words that appear most frequently in all documents of the collection; these words would become my stop words list.

5. Please describe how you process phrase and boolean queries and any optimizations you added for faster processing.

I used a query factory structure for generating phrase and boolean objects. However, my processing pipeline is different for each query. In the case of phrase, I first did all

of my remove characters, lowercase, and stopword matching prior to matching the query to my inverted index. However, in the boolean case, my pipeline was in a different order; I first sent the un-processing query to the provided boolparser.py file and generated an abstract syntax tree (AST). Then, within the recursive function, I did my processing (e.g. lowercase, stemming, and matching) for each of the terms in my query.